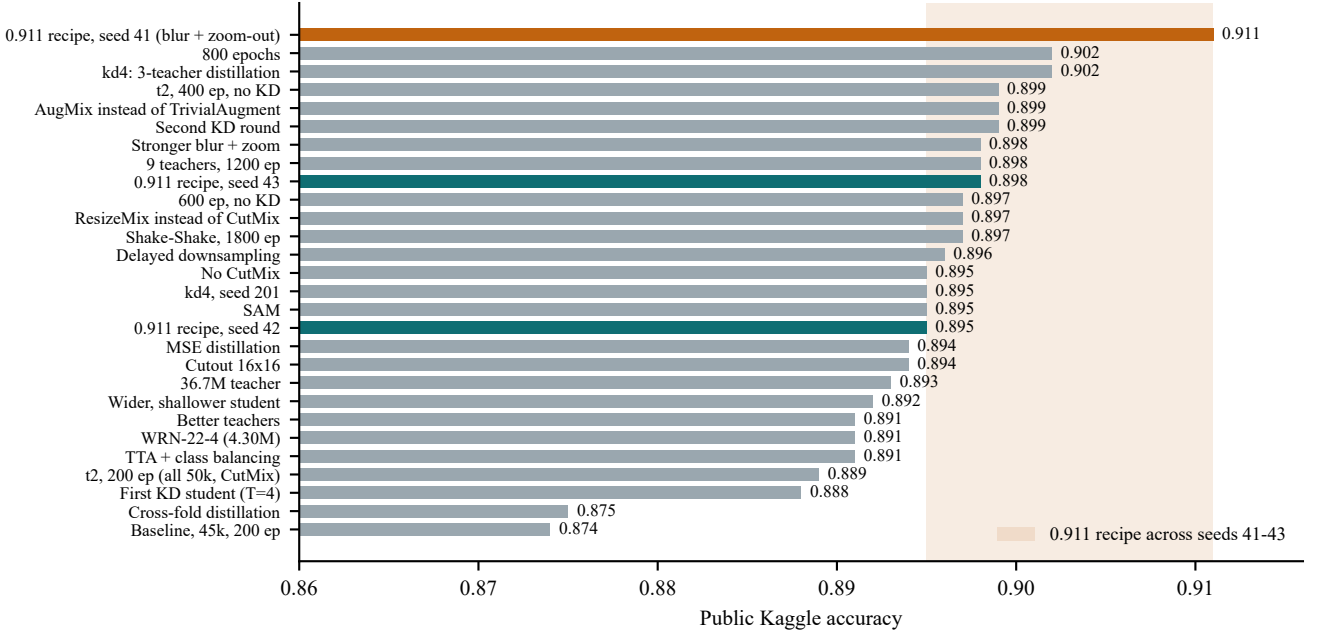
0.911 recipe, seed 41 (blur + zoom-out)
800 epochs
kd4: 3-teacher distillation
t2, 400 ep, no KD
AugMix instead of TrivialAugment
Second KD round
Stronger blur + zoom
9 teachers, 1200 ep
0.911 recipe, seed 43
600 ep, no KD
ResizeMix instead of CutMix
Shake-Shake, 1800 ep
Delayed downsampling
No CutMix
kd4, seed 201
SAM
0.911 recipe, seed 42
MSE distillation
Cutout 16x16
36.7M teacher
Wider, shallower student
Better teachers
WRN-22-4 (4.30M)
TTA + class balancing
t2, 200 ep (all 50k, CutMix)
First KD student (T=4)
Cross-fold distillation
Baseline, 45k, 200 ep
0.911
0.902
0.902
0.899
0.899
0.899
0.898
0.898
0.898
0.897
0.897
0.897
0.896
0.895
0.895
0.895
0.895
0.894
0.894
0.893
0.892
0.891
0.891
0.891
0.889
0.888
0.875
0.874
0.911 recipe across seeds 41-43
0.86
0.87
0.88
0.89
0.90
0.91
Public Kaggle accuracy