Deep Learning Project 1

Chetan Kodeboyina (ck3399) Shwetanshu Raj (sr8250) Ninad Rade (nr3263)
Team *Jugaad Net*
New York University, ECE-GY 7123 Deep Learning, Project 1

Abstract

We train a residual network with fewer than 5M parameters from scratch on the 50,000 CIFAR-10 training images and predict a hidden 2,000-image test set drawn from a newer, shifted collection of images. Over about 30 training runs and 33 leaderboard submissions, the public score rose from 0.874 to 0.911. The final model is a 4.70M-parameter ResNet trained for 400 epochs with strong augmentation, blur and zoom-out to target the shift, and distillation from four same-size teachers. This report describes the ideas we tested and what we learned from them. First, most of the gain came from using all the data, the full parameter budget and enough epochs. Second, held-out accuracy on training images was close to 98.5% for every variant and did not predict the leaderboard. Third, the public score was too noisy to compare our variants. Rerunning our best recipe with two new seeds scored 0.895 and 0.898, and on the private half of the test set our selected 0.911 submission scored 0.876, below seven submissions we did not choose.

1 Introduction

The task is CIFAR-10 classification (Krizhevsky, 2009) under three rules: a single model, at most 5M trainable parameters, trained from scratch with no external data. Models are scored on 2,000 hidden images, 200 per class. The public leaderboard uses about half of them, so each image is worth 0.001 of accuracy.

What makes the task hard is that the hidden images come from a newer collection than the training images. Recht et al. (2018) showed that such a collection, gathered the same way, costs CIFAR-10 models several points of accuracy for reasons that are not fully understood. Our models score about 98.5% on held-out training images but around 0.90 on the leaderboard. So most of our work was about making a small model hold up on slightly different photos.

All local measurements in this report use a fixed 5,000-image held-out split of the training set, on models trained with the other 45,000 images. The hidden test images were never labelled by hand for model selection.

Main takeaways.

- Most of the gain came from the basics. All 50k images, a network deep enough to use the parameter budget, CutMix and 400 epochs took the score from 0.874 to 0.899.
- Blur and zoom-out augmentation, aimed at the shift, produced our best public score, 0.911.
- Longer training, larger teachers and stronger augmentation all scored lower.
- Held-out accuracy did not track the leaderboard, and seed variance (± 0.01) was as large as the effects we were measuring.
- Choosing the final submissions by the highest public score did not work well, because public score did not predict private score.

2 The Recipe in Brief

We kept the model family fixed and spent our effort on data, training and evaluation. The network is a standard basic-block ResNet (He et al., 2016): a 64-channel stem and three stages of 64, 128 and 256 channels with 4, 4 and 3 blocks, global average pooling and a linear layer (4,697,162 parameters, Table 1). Training uses SGD with cosine learning-rate decay (Loshchilov and Hutter, 2017), weight averaging (EMA), label smoothing (Szegedy et al., 2016) and four kinds of augmentation:

- **General:** random crops and flips, TrivialAugment (Müller and Hutter, 2021) and small random erasing (Zhong et al., 2020).
- **Mixing:** CutMix (Yun et al., 2019), which pastes a patch of one image into another and mixes the labels.
- **Shift-targeted:** Gaussian blur on 20% of images and zoom-out on 30%, which shrinks the object inside the frame.
- **Distillation:** the model learns to match the averaged predictions of four same-size teachers, following the "patient and consistent" setup of Beyer et al. (2022).

At test time we average predictions over the image, its mirror image and small shifts.

3 How the Score Moved

Figure 1 shows the five steps that raised the public score.

Use everything we are allowed (0.874 $\rightarrow$ 0.889). Our baseline held out 5k images for validation and used a (3,3,3) network. Training on all 50k images, adding blocks up to the 5M budget and adding CutMix gave the largest single jump.

Stage	Channels	Blocks	Feature map
Stem	64	–	32×32
Stage 1	64	4	32×32
Stage 2	128	4	16×16
Stage 3	256	3	8×8
Pool + linear	10	–	1×1
Parameters: **4,697,162** of 5,000,000 allowed			

Table 1: Final architecture: basic residual blocks with 3×3 kernels.

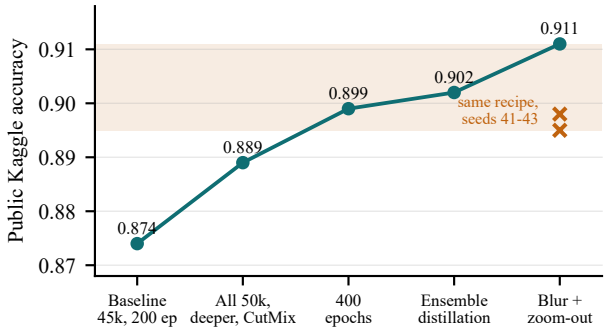


Figure 1: Best public score after each milestone. Crosses: the final recipe rerun with seeds 42 and 43. The shaded band is the spread of that one recipe across three seeds.

Before trying more complex ideas, it is worth using all the data and all the capacity the rules allow.

Train longer (0.889 → 0.899). Doubling training from 200 to 400 epochs added a full point. Going further did not: 600, 800 and 1,200 epochs all scored at or below 400. Figure 2 shows where the gain comes from. Held-out accuracy rises from 95.7% at epoch 200 to 98.5% at epoch 400, almost all of it while the cosine schedule lowers the learning rate, and the training loss falls in step. A longer schedule only spreads this late improvement over more epochs.

Learn from an ensemble (0.899 → 0.902). An ensemble of models is more accurate than any one of them, but the rules allow only a single model. Distillation lets one model learn the ensemble's averaged predictions. The version that worked used soft targets only (no hard labels) and showed the teachers the same augmented image as the student (Beyer et al., 2022). Earlier variants scored lower, including a higher temperature (0.888) and teachers that never saw the image they labelled (0.875).

Target the shift (0.902 → 0.911). We measured how our models fail on degraded images (Section 4) and found two weak spots: blur and small objects. Adding blur and zoom-out to training produced our best score.

Check the result (0.911 → 0.895). At the end we reran the exact 0.911 recipe with two new seeds. They scored 0.895 and 0.898. This changed how we interpret the other results (Section 8).

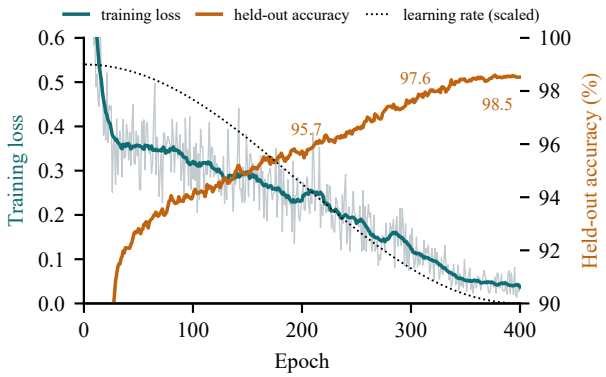


Figure 2: Training loss (smoothed; raw values in grey) and held-out accuracy of the final recipe trained on 45k images, evaluated each epoch on the 5k images it never trained on (EMA weights, image + mirror). Most of the held-out gain comes in the second half, as the learning rate decays.

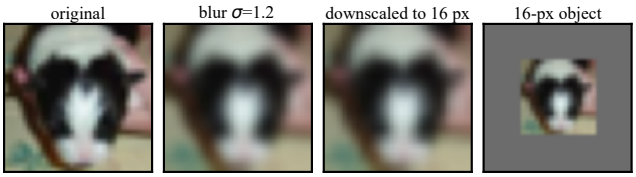


Figure 3: The distortion tests applied to a held-out training image.

4 Probing Robustness and Errors

Because newer photos can differ in sharpness and framing, we asked how the model copes with degraded versions of images it has never seen. All numbers in this section use the 5,000 held-out training images and models trained on the other 45,000.

Distortion tests. We applied controlled distortions (Figure 3) and measured accuracy. Without the shift-targeted augmentation, the distilled model dropped from 98.4% to 83.7% under blur (σ=1.2) and to 79.1% when the object filled only 16 pixels. Cats and dogs suffered most. Adding blur and zoom-out to training raised these to 86.9% and 81.9% while leaving clean accuracy unchanged (Figure 4). On the leaderboard the same change was worth 0.009, from 0.902 to 0.911, though Section 8 shows that this gap is within seed noise.

Error analysis. The final recipe makes 68 errors on the 5,000 held-out images (98.6% with test-time averaging). Figure 5 counts the most common confusions and Figure 6 shows the most confident ones. Cat and dog account for 25 of the 68. The rest are spread thinly across many pairs, mostly small or partly hidden animals and unusual viewpoints. Many of the confident mistakes are hard even for a person at 32×32 pixels.

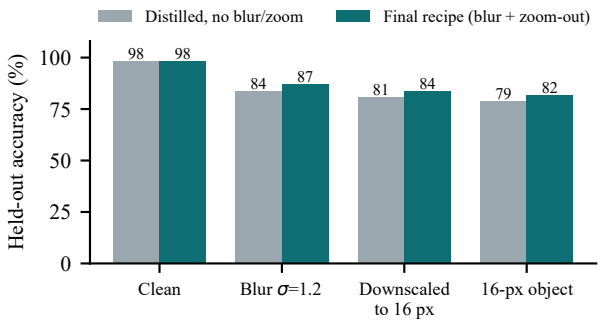


Figure 4: Held-out accuracy (single view) of the distilled model before and after adding blur and zoom-out augmentation. Robustness improves without any cost on clean images.



Figure 5: Most frequent confusions of the final recipe on the held-out images. Cat and dog are the largest group.

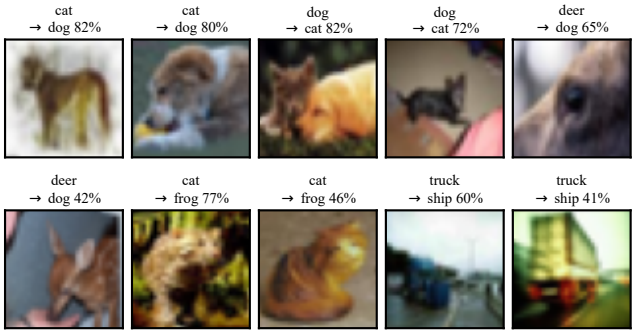


Figure 6: The most confident held-out errors for the main confusion pairs (true class → prediction, confidence).

Model	Held-out (45k)	Public
Baseline (3,3,3), 200 ep	97.22	0.874
Distilled, 3 teachers	98.56	0.902
36.7M teacher / its student	97.84 / –	– / 0.893
Final (blur + zoom-out)	**98.64**	**0.911**
same, seed 42	–	0.895
same, seed 43	–	0.898
+ grayscale	98.70	no gain
800 epochs	–	0.902
no CutMix	–	0.895
ResizeMix	–	0.897
Shake-Shake 2×40d, 1800 ep	–	0.897

Table 2: Held-out accuracy (%, 5k images, models trained on 45k, test-time averaging) and public score of the matching model trained on all 50k. "–": no 45k version was trained.

The limit of held-out data. Held-out images come from the same collection as the training images, so they cannot show what is different about the hidden set. Every distilled variant we trained on 45k images landed between 98.5% and 98.7% held-out, while the same variants differed by up to 0.02 on the leaderboard.

5 Ideas Pursued

Besides the milestones above, we followed six lines of work. Each started from a specific hypothesis.

1. **More capacity per parameter.** A wider, shallower network and a Shake-Shake network (Gastaldi, 2017) might use the 5M budget better than a standard ResNet.
2. **Better optimization.** Sharpness-aware minimization (Foret et al., 2021) seeks flatter minima, which some work suggests transfer better to new data.
3. **Better teachers.** If distillation helps, a stronger teacher should help more: we tried a 36.7M-parameter teacher, nine teachers, and teachers that were themselves distilled.
4. **Other augmentations.** AugMix (Hendrycks et al., 2020), Cutout (DeVries and Taylor, 2017), grayscale and stronger blur and zoom each target a different kind of robustness.
5. **Different mixing.** ResizeMix (Qin et al., 2020) pastes a whole shrunken image instead of a patch, so the pasted object is always complete; zoom-in crops target close-up views.
6. **Teacher-free robustness.** Teachers trained without blur might hold the student back, so we trained blur-augmented models without teachers and distilled from those.

None of these improved the public score. The next three sections explain why we think so.

6 Results

Table 2 summarizes the main models and Figure 7 shows every submission. Held-out accuracy, where we have it, barely moves; the leaderboard spans 0.874 to 0.911.

7 What We Expected to Work but Didn't

Table 3 lists the ideas we were most confident about. We see three patterns.

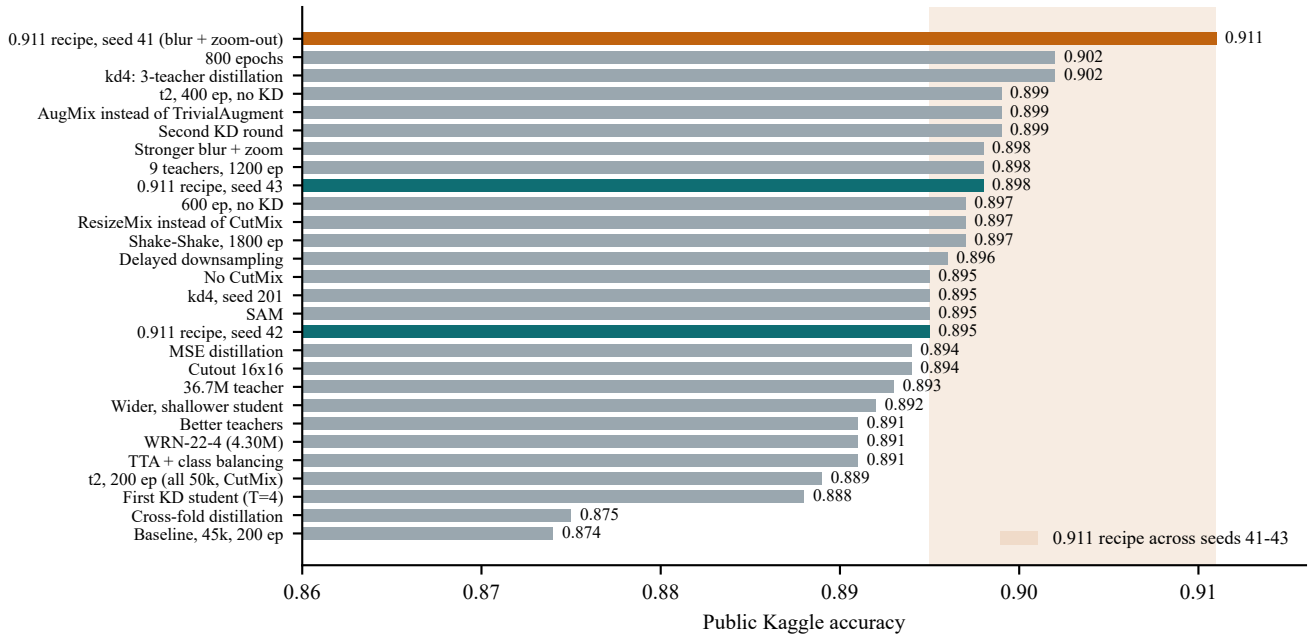


Figure 7: Public score of every scored submission. The shaded band is the spread of the final recipe across seeds 41–43 (0.895 to 0.911); most variants fall inside it. Not shown: grayscale augmentation (no gain, score not recorded) and seven early submissions without a recorded description.

Idea	Public	Outcome
600 / 800 / 1200 epochs	.897/.902/.898	Never better than 400
36.7M or better teachers	.893/.891	Stronger teacher, weaker student
Sharpness-aware min.	.895	No gain, twice the cost
Stronger blur + zoom	.898	Too much distortion
Shake-Shake, 1800 ep	.897	Strong on paper, no gain
ResizeMix	.897	No gain over CutMix
Grayscale	no gain	+0.06 held-out only
Other distillation variants	.875–.899	None beat the simple version
Balancing class counts	.891	Below plain averaging (0.899)

Table 3: Ideas we expected to help, with public scores.

Pushing working ideas further did not help. Longer training, more teachers and stronger augmentation had each helped in a mild form. Pushed further, they did as well on the training distribution but never better on the hidden set.

Stronger teachers gave weaker students. A 36.7M-parameter teacher reached 97.84% held-out after only 200 epochs, but its student scored 0.893, below students of our same-size teachers. One possible reason is that a small student finds it harder to match a much larger, very different network.

Targeted fixes were too small to see. ResizeMix, zoom-in and teacher-free distillation each aimed at a narrow kind of error. Even if they worked, they would change only a handful of the 1,000 public images, below the noise floor of one submission.

8 How Much Was Noise

Seeds. Rerunning the final recipe with seeds 42 and 43 scored 0.895 and 0.898, against 0.911 for seed 41 (Figure 8). An earlier recipe varied by 0.005 across three seeds. Two seeds of the same recipe disagree on 72–80 of the 2,000 hidden images. That is more than our 800-epoch run (58) or ResizeMix run (68) changed relative to 0.911. So most of the "failed" variants in Table 3 were compared against a lucky draw, and some may be as good as the recipe's typical result.

Public versus private. After the competition, Kaggle revealed each submission's score on the other half of the hidden set. We had chosen our two final submissions by public score: 0.911 and 0.902, the two highest. On the private half both scored 0.876, which ranked us 22nd (16th on public). Seven submissions we did not choose scored higher, the best at 0.882 (Figure 9). Across our 31 submissions from the full training set, the rank correlation between public and private score is only 0.18 ($p = 0.34$). The private half is harder for everyone, about 0.024 lower on average, but the order of our models within it is close to random. This is known as the winner's curse. The highest public score is partly the luckiest one, so choosing by it selects luck as well as quality. The seed reruns agree with this. On private, seeds 41, 42 and 43

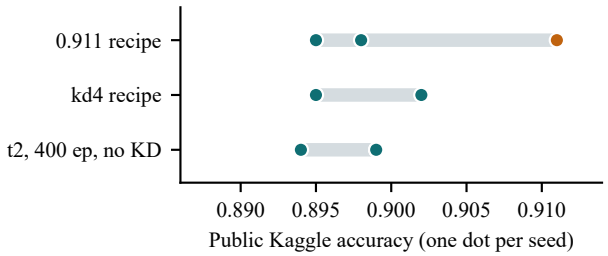


Figure 8: Public score of the same recipe with different seeds. The final recipe spans 0.016, the size of most effects we tried to measure.

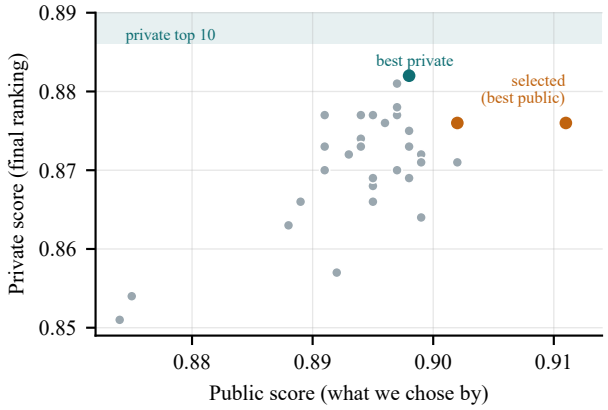


Figure 9: Public against private score for all 33 submissions. The two we selected had the best public scores but only average private scores. Shaded: the private score needed for the top 10.

scored 0.876, 0.877 and 0.869.

Held-out accuracy. Our held-out split is saturated. The distilled variants we trained on 45k images differ by at most 0.15 points, about the noise of 5,000 images, while their full-data versions differ by up to 0.02 on the leaderboard. For example, grayscale augmentation was 0.06 points better held-out but gave no gain on the leaderboard.

Hand labels. We also produced two independent sets of labels for the hidden images to use as a local check. The two sets disagreed on 503 of 2,000 images, and neither ranked our submissions in the same order as the leaderboard (rank correlation 0.18–0.42, not significant). We did not use them to choose models.

9 What Surprised Us

- **Our best score was largely luck.** The same recipe scored 0.911, 0.895 and 0.898 on public, and 0.876, 0.877 and 0.869 on private.
- **Our best public submissions were not our best private ones.** Seven unselected submissions beat them on the private half.
- **Robustness gains were real but small on the leaderboard.** Blur and zoom-out added 3 points under blur held-out and at most about one point on the leaderboard.
- **Bigger was worse.** A teacher seven times larger, and training three times longer, both gave lower scores than the simple recipe.
- **Cat and dog stayed the hardest pair.** They account for over a third of the final recipe's held-out errors, more than any augmentation changed.

10 Lessons Learned

1. **Measure the noise first.** One rerun with a new seed, done early, would have shown that differences under 0.02 cannot be read from a single submission. We would have tested fewer ideas, each with three seeds.
2. **Do not pick final submissions by the single highest public score.** Choose the recipe with the best average over seeds, and hedge the second pick with a model that is strong for a different reason.
3. **Get the basics right first.** Data, depth and training length gave 0.025 of our 0.037 total gain. None of the more complex ideas added as much.
4. **A validation set from the training distribution cannot see a shift.** Held-out accuracy was saturated near 98.5% and did not rank our models the way the leaderboard did. Testing robustness directly, with distortions, was more informative.
5. **Look at the errors early.** The error grids showed where the model struggled (cats and dogs, small and partly hidden animals) more clearly than any single metric.
6. **Robustness to one distortion does not guarantee robustness to a shift.** Blur and zoom-out clearly helped on degraded images, but the hidden set differs in ways we could not measure locally.

11 Conclusion and Future Work

A 4.7M-parameter ResNet trained from scratch reaches 98.6% on held-out training images, 0.895–0.911 on the public half of the shifted hidden set and 0.869–0.882 on the private half. Most of that came from standard training practice. We could not measure the effect of our more targeted ideas, because the leaderboard was too noisy and our held-out split was too easy to separate them. With more time we would (1) run three seeds per idea and compare averages, (2) train the teachers with the same blur and zoom-out augmentation so distillation passes on robustness, and (3) average the seeds of the final recipe into one distilled model, which would reduce the seed variance we measured.

Code. The training script, run scripts, logs and full experiment log are included in the code archive submitted with this report.

References

Beyer, L.; Zhai, X.; Royer, A.; Markeeva, L.; Anil, R.; and Kolesnikov, A. 2022. Knowledge distillation: A good teacher is patient and consistent. In *CVPR*.

DeVries, T.; and Taylor, G. W. 2017. Improved regularization of convolutional neural networks with Cutout. *arXiv:1708.04552*.

Foret, P.; Kleiner, A.; Mobahi, H.; and Neyshabur, B. 2021. Sharpness-aware minimization for efficiently improving generalization. In *ICLR*.

Gastaldi, X. 2017. Shake-Shake regularization. *arXiv:1705.07485*.

He, K.; Zhang, X.; Ren, S.; and Sun, J. 2016. Deep residual learning for image recognition. In *CVPR*.

Hendrycks, D.; Mu, N.; Cubuk, E. D.; Zoph, B.; Gilmer, J.; and Lakshminarayanan, B. 2020. AugMix. In *ICLR*.

Hinton, G.; Vinyals, O.; and Dean, J. 2015. Distilling the knowledge in a neural network. *arXiv:1503.02531*.

Krizhevsky, A. 2009. Learning multiple layers of features from tiny images. Technical report, University of Toronto.

Loshchilov, I.; and Hutter, F. 2017. SGDR: Stochastic gradient descent with warm restarts. In *ICLR*.

Müller, S. G.; and Hutter, F. 2021. TrivialAugment: Tuning-free yet state-of-the-art data augmentation. In *ICCV*.

Qin, J.; Fang, J.; Zhang, Q.; Liu, W.; Wang, X.; and Wang, X. 2020. ResizeMix: Mixing data with preserved object information and true labels. *arXiv:2012.11101*.

Recht, B.; Roelofs, R.; Schmidt, L.; and Shankar, V. 2018. Do CIFAR-10 classifiers generalize to CIFAR-10? *arXiv:1806.00451*.

Szegedy, C.; Vanhoucke, V.; Ioffe, S.; Shlens, J.; and Wojna, Z. 2016. Rethinking the Inception architecture for computer vision. In *CVPR*.

Yun, S.; Han, D.; Oh, S. J.; Chun, S.; Choe, J.; and Yoo, Y. 2019. CutMix: Regularization strategy to train strong classifiers with localizable features. In *ICCV*.

Zhong, Z.; Zheng, L.; Kang, G.; Li, S.; and Yang, Y. 2020. Random erasing data augmentation. In *AAAI*.